

▼ Child Seat Localization

Author: Fernando Kelvin da Silva Soares

Notebook Versio: V0.0

▼ Introduction

In this notebook we will explore the child seat localization problem. The goal here, is to say in wich seat a child/infant seat is locates in the back part of a veicle. It was used as reference the SVIRO sunthetic dataset and some papers [1][2] wrote by the sviro team, available in [the SVIRO website](#), as weel as, other cientific papers and websites to implement a image classifier for each back seat position.

Differently from the implementations presented by the SVIRO team on their papers, wich are focused on train the networks in one vehicle and see how they perform in unknown vchiles. In this work was decided to use more then one veichle for trainin and check the accuracy in unknown vehicles, to explore if there is a significant reseult compared to the single vehicle training.

The porpouse of this project is to give a flexible and fast framework to explore the child seat problem understanding the limitations and possibilities involved.

▼ Copyright The SVIRO Authors

The usage of the SVIRO dataset is subjectet to the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. All all the credits must be given to the authors as follow.

[1] Steve Dias Da Cruz, Oliver Wasenmüller, Hans-Peter Beise, Thomas Stifter, & Didier Stricker (2020). SVIRO: Synthetic Vehicle Interior Rear Seat Occupancy Dataset and Benchmark. In IEEE Winter Conference on Applications of Computer Vision (WACV).

[2] Steve Dias Da Cruz, Bertram Taetz, Oliver Wasenmüller, Thomas Stifter, & Didier Stricker (2021). Autoencoder Based Inter-Vehicle Generalization for In-Cabin Occupant Classification. In IEEE Intelligent Vehicles Symposium (IV).

Attribution-NonCommercial-ShareAlike 4.0 International

▼ Copyright 2019 The TensorFlow Authors.

Many function on this notebook ware based on the tutorial https://www.tensorflow.org/tutorials/images/transfer_learning, and all the rights are reseverved to the authors. Please do not share this notebook witout given proper credits.

Licensed under the Apache License, Version 2.0 (the "License");

MIT License

▼ Imports

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import random
4 import mathlib
```

```

4 import pathlib
5 import seaborn as sns
6 from os import listdir
7 import os.path as path
8 import os
9 import PIL
10 import tensorflow as tf
11
12 from tensorflow import keras
13 from tensorflow.keras import layers
14 from tensorflow.keras.models import Sequential
15 from tensorflow.keras.preprocessing import image_dataset_from_directory
16 from tensorflow.keras.layers.experimental.preprocessing import RandomRotation
17 from tensorflow.keras.layers.experimental.preprocessing import RandomFlip
18 from tensorflow.keras.layers.experimental.preprocessing import RandomZoom
19 from tensorflow.keras.layers.experimental.preprocessing import RandomContrast
20
21 #Check the tensorflow version, the recommended for this notebook is the 2.5.0.
22 print(tf.__version__)

```

2.5.0

If you are using Google Colab, it is recommended to use the GPU for fast training. To activate the GPU, go to **"Runtime->Change runtime type"** and select GPU at in the **"Hardware Accelerator"** dropdown menu.

You can check the allocated GPU model running the following command.

```
1 !nvidia-smi
```

```

Wed Jun 16 03:00:30 2021
+-----+
| NVIDIA-SMI 465.27      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+-----+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M. |
+-----+-----+-----+-----+-----+-----+
|    0  Tesla P100-PCIE...    Off   | 00000000:00:04.0 Off |             0        |
| N/A   34C    P0      26W / 250W |  0MiB / 16280MiB |      0%      Default |
+-----+-----+-----+-----+-----+-----+
|
+-----+-----+-----+-----+-----+
| Processes: |
| GPU   GI    CI          PID    Type    Process name                        GPU Memory |
|          ID    ID                                   Usage   |
+-----+-----+-----+-----+-----+
| No running processes found |
+-----+

```

▼ Conecto to your Drive repository

If you want to use your Google Drive repository to load the train dataset, run the following session, log-in into your account, copy the verification code and paste in the output entry.

```

1 from google.colab import drive
2
3 drive.mount('/content/gdrive')

```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive

▼ Global Definitions

Here we have some constants that are used in all the sections of the notebook.

Set up the MAIN PATH to you repository main folder containing, the models and datasets.

The IMAGE_SHAPE will depend on the network you are using. In this case we will train a [EfficientNet stadard implementation](#) network that accepts different input sizes, so can use the single seat image size of 250x550.

In this notebook we can explore not only the SVIRO classes but also a summarized version of the existent clsses, where all the child/infant seat with or without child are cosidered as the same class.

```

1 MAIN_PATH = '/content/gdrive/MyDrive/Colab_Notebooks/Child_Seat_Localization/classifier'
2
3 IMAGE_HIGHT = 258
4 IMAGE_WIDTH = 258
5 IMAGE_SHAPE = (IMAGE_HIGHT, IMAGE_WIDTH, 3) #The image shape may vary according to you network model
6
7 # SVIRO Classes
8 CLASSES = {0: "Empty seat",
9             1: "Infant in infant seat",
10            2: "Child on child seat",
11            3: "Adult passenger",
12            4: "Everyday object",
13            5: "Infant seat without baby",
14            6: "Child seat without child"}
15

```

▼ Preparing the Dataset

As we are implementing a sigle seat image cassifier the training dataset was created using the single seat grayscale images fom the SVIRO dataset.

In the original dataset we have 10 different vehicles, that we can arrange in 3 classes: Small Vhicle (two door), Regular Vehicle and Big vehicle (SUV and Truck). So, to create the train/validation dataset one vehicle from each one of this classes was chosen, considereng a good variability of interior styles. Are they: Renault **Zoe**, Toyota **Hilux** and Tesla **Model3**.

From the original train folder of each one, 1800 images were taken for each of the seven classes, totalizing 3150 images on the final train/validation dataset.

It was used a 80:20 split for training and validation data.

```

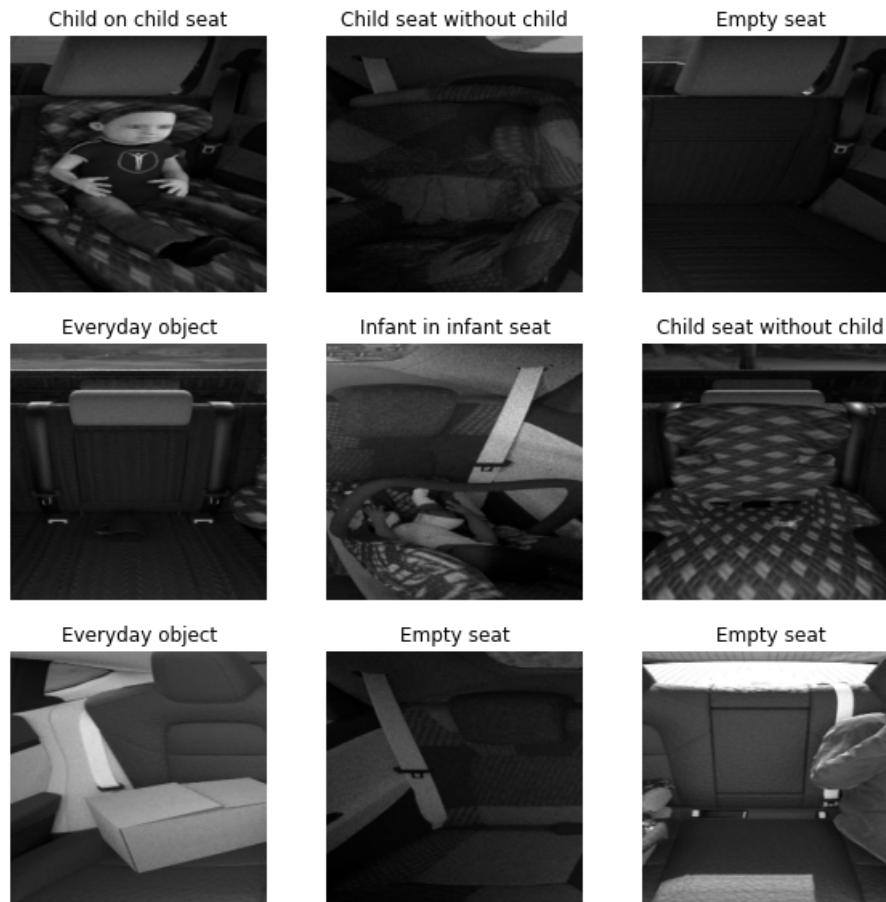
1 train_dataset_path = '{} /datasets/train/sviro_classes'.format(MAIN_PATH)
2
3 # From the paper of efficient det they use a 258x258 resolution for the b7 model
4 BATCH_SIZE = 64
5 VALIDATION_SPLIT = 0.2
6
7 # crate a train and validation dataset with a 80:20 split
8 train_dataset = image_dataset_from_directory(train_dataset_path,
9                                              validation_split=VALIDATION_SPLIT,
10                                             subset="training",
11                                             seed=123,
12                                             image_size=(IMAGE_HIGHT, IMAGE_WIDTH),
13                                             color_mode="rgb",
14                                             batch_size=BATCH_SIZE,
15                                             smart_resize=False)
16
17 valid_dataset = image_dataset_from_directory(train_dataset_path,
18                                              validation_split=VALIDATION_SPLIT,
19                                              subset="validation",
20                                              seed=123,
21                                              image_size=(IMAGE_HIGHT, IMAGE_WIDTH),
22                                              color_mode="rgb",
23                                              batch_size=BATCH_SIZE,
24                                              smart_resize=False)

```

Found 3150 files belonging to 7 classes.
 Using 2520 files for training.
 Found 3150 files belonging to 7 classes.
 Using 630 files for validation.

Looking at some examples from the train dataset.

```
1 plt.figure(figsize=(10, 10))
2 for images, labels in train_dataset.take(1):
3     for i in range(9):
4         ax = plt.subplot(3, 3, i + 1)
5         plt.imshow(images[i].numpy().astype("uint8"))
6         plt.title(CLASSES[int(labels[i])])
7         plt.axis("off")
```



Here we use AUTOTUNE prefetch to optimise the traing performance.

```
1 AUTOTUNE = tf.data.experimental.AUTOTUNE
2
3 train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
4 validation_dataset = valid_dataset.prefetch(buffer_size=AUTOTUNE)
```

▼ Training the Model

▼ Create the model based on a pre-trained model

For a faster implemantation and training, it was chosen to use the pre-trained models, on the imagenet dataset, from [Keras Aplication](#) module.

As said before, for this implementation the EfficientNet B7 was chosen, due to its good performance on the imagenet dataset and faster inference time compared to the best existing neural networks, as we can see in the [paperswithcode.com Image Classification on ImageNet benchmark](https://paperswithcode.com/Image-Classification-on-ImageNet-benchmark).

An additional data augmentation layer was placed inside the model during training to avoid overfitting.

```

1 number_of_classes = len(CLASSES)
2
3 data_augmentation = keras.Sequential([
4     RandomRotation(0.2, input_shape=IMAGE_SHAPE),
5     RandomFlip(mode="horizontal"),
6     RandomContrast(factor=0.2)
7 ])
8
9 pretrained_model = tf.keras.applications.EfficientNetB7(input_shape=IMAGE_SHAPE,
10                                                         include_top=False,
11                                                         weights="imagenet")
12 # pretrained_model.summary()
13
14 #Freeze the original convolutional weights from the pre-trained model
15 pretrained_model.trainable=False
16
17 model = tf.keras.Sequential([
18     data_augmentation,
19     pretrained_model,
20     layers.GlobalAveragePooling2D(),
21     layers.Dense(number_of_classes, activation="softmax")
22 ])
23
24 model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 258, 258, 3)	0
efficientnetb7 (Functional)	(None, 9, 9, 2560)	64097687
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2560)	0
dense (Dense)	(None, 7)	17927
Total params: 64,115,614		
Trainable params: 17,927		
Non-trainable params: 64,097,687		

▼ Compiling the model

For optimizing the model the Adam optimizer was used, for being easier to use and have a good performance on most problems.

```

1 model.compile(optimizer='adam',
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
3               metrics=['accuracy'])

```

▼ Create checkpoints and train the model

As only the last layers (classification layers) will be trained, the model is fitted for only 25 epochs. And a checkpoint will be saved each time we get a better accuracy result on the validation dataset.

```

1 checkpoint_path = "{}models/EfficientNet_b7_258_258_RFC/".format(MAIN_PATH)
2
3 # Create a callback that saves the model's weights
4 cp_callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path,
5                                                    save_weights_only=True,
6                                                    verbose=1,
7                                                    monitor='val_accuracy',
8                                                    mode='max',
9                                                    save_best_only=True)
10
11
12 # If you want to train from the last chachpoint uncomment the section below
13 # latest_checkpoint = tf.train.latest_checkpoint(checkpoint_path)
14 # print(latest_checkpoint)
15
16 # if not latest==None:
17 #     model.load_weights(latest)
18
19 # Train the model with the new callback
20 epochs = 25
21 history = model.fit(train_dataset,
22                     validation_data=valid_dataset,
23                     epochs=epochs,
24                     callbacks=[cp_callback]) # Pass callback to training

```

Epoch 1/25
40/40 [=====] - 54s 900ms/step - loss: 1.1647 - accuracy: 0.6865 - val_loss: 0.5

Epoch 00001: val_accuracy improved from -inf to 0.89365, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00001.h5
Epoch 2/25
40/40 [=====] - 32s 772ms/step - loss: 0.5463 - accuracy: 0.8829 - val_loss: 0.3

Epoch 00002: val_accuracy improved from 0.89365 to 0.93016, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00002.h5
Epoch 3/25
40/40 [=====] - 32s 771ms/step - loss: 0.3933 - accuracy: 0.9063 - val_loss: 0.2

Epoch 00003: val_accuracy improved from 0.93016 to 0.95238, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00003.h5
Epoch 4/25
40/40 [=====] - 32s 772ms/step - loss: 0.3288 - accuracy: 0.9274 - val_loss: 0.2

Epoch 00004: val_accuracy improved from 0.95238 to 0.96032, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00004.h5
Epoch 5/25
40/40 [=====] - 33s 773ms/step - loss: 0.2793 - accuracy: 0.9333 - val_loss: 0.2

Epoch 00005: val_accuracy did not improve from 0.96032
Epoch 6/25
40/40 [=====] - 32s 772ms/step - loss: 0.2618 - accuracy: 0.9317 - val_loss: 0.1

Epoch 00006: val_accuracy improved from 0.96032 to 0.96349, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00006.h5
Epoch 7/25
40/40 [=====] - 32s 774ms/step - loss: 0.2168 - accuracy: 0.9492 - val_loss: 0.1

Epoch 00007: val_accuracy did not improve from 0.96349
Epoch 8/25
40/40 [=====] - 32s 772ms/step - loss: 0.2139 - accuracy: 0.9433 - val_loss: 0.1

Epoch 00008: val_accuracy improved from 0.96349 to 0.96984, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00008.h5
Epoch 9/25
40/40 [=====] - 32s 773ms/step - loss: 0.1911 - accuracy: 0.9520 - val_loss: 0.1

Epoch 00009: val_accuracy improved from 0.96984 to 0.97143, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00009.h5
Epoch 10/25
40/40 [=====] - 32s 773ms/step - loss: 0.1853 - accuracy: 0.9548 - val_loss: 0.1

Epoch 00010: val_accuracy did not improve from 0.97143
Epoch 11/25
40/40 [=====] - 32s 773ms/step - loss: 0.1712 - accuracy: 0.9579 - val_loss: 0.1

Epoch 00011: val_accuracy did not improve from 0.97143
Epoch 12/25
40/40 [=====] - 32s 772ms/step - loss: 0.1662 - accuracy: 0.9544 - val_loss: 0.1

Epoch 00012: val_accuracy improved from 0.97143 to 0.97302, saving model to /content/gdrive/MyDrive/Colab_Notebooks/child_seat_localization/child_seat_localization_20210616_140000_00012.h5

```

Epoch 13/25
40/40 [=====] - 32s 773ms/step - loss: 0.1503 - accuracy: 0.9611 - val_loss: 0.1

Epoch 00013: val_accuracy did not improve from 0.97302
Epoch 14/25
40/40 [=====] - 32s 772ms/step - loss: 0.1549 - accuracy: 0.9615 - val_loss: 0.1

Epoch 00014: val_accuracy improved from 0.97302 to 0.97460, saving model to /content/gdrive/MyDrive/Colab
Epoch 15/25
40/40 [=====] - 32s 772ms/step - loss: 0.1395 - accuracy: 0.9639 - val_loss: 0.1

```

▼ Visualize training results


Create plots of loss and accuracy on the training and validation sets.

```

1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')

```

Text(0.5, 1.0, 'Training and Validation Loss')



▼ Saving the model

We can save the entire model for use in future applications.

```
1 save_model_path = '{}/models/EfficientNet_b7_258_258_RFC/model'.format(MAIN_PATH)
2 model.save(save_model_path)
```

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/utils/generic_utils.py:497: CustomMaskWarning
category=CustomMaskWarning)
INFO:tensorflow:Assets written to: /content/gdrive/MyDrive/Colab_Notebooks/Child_Seat_Localization/classifi
```

▼ Testing the model performance

▼ Load the model

It is possible to load different models to evaluate the performance. Pay attention on the IMAGE_HEIGHT and IMAGE_WIDTH to match the model input size.

```
1 load_model_path = '{}/models/EfficientNet_b7_258_258_RFC/model'.format(MAIN_PATH)
2
3 model = tf.keras.models.load_model(load_model_path)
```

```
WARNING:absl:Importing a function (__inference_block4b_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block3a_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block1a_activation_layer_call_and_return_conditional_losses)

WARNING:absl:Importing a function (__inference_block4b_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5j_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5c_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6i_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block3e_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2a_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2g_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2b_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2f_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6j_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6l_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block3b_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4a_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5c_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6j_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6b_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5f_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4c_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block1d_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5g_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block1c_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block7d_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6l_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5a_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6d_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6f_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4a_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block1d_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4g_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4i_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2e_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6l_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4c_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5h_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6j_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6a_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5g_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5f_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5d_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4j_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block3g_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6j_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4g_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block7b_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block1c_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5i_se_reduce_layer_call_and_return_conditional_losses)
```



```

WARNING:absl:Importing a function (__inference_block5i_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block5g_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6k_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2b_expand_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6g_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block3c_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block4h_se_reduce_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block6i_activation_layer_call_and_return_conditional_losses)
WARNING:absl:Importing a function (__inference_block2a_expand_activation_layer_call_and_return_conditional_losses)

```

▼ Evaluate the model performance and plot the confusion matrix

The model is evaluated on the single seat grayscale test images from each unknown vehicle (aclass, escape, gsf, i3, tiguan, tucson and x5) and plot the confusion matrix.


```

1 # path to the original SVIRO dataset containing the grayscale single seat images
2 # of each vehicle
3 original_dataset_path = '{} / datasets / test / grayscale_single_seat'.format(MAIN_PATH)
4
5 vehicles_to_test = ['aclass', 'escape', 'gsf', 'i3', 'tiguan', 'tucson', 'x5']
6
7 cars_accuracy = {}
8 for vehicle in vehicles_to_test:
9
10     test_dataset_path = original_dataset_path + \
11         "{} / test_with_labels / grayscale".format(vehicle)
12
13     test_dataset = image_dataset_from_directory(test_dataset_path,
14                                                 validation_split=None,
15                                                 subset=None,
16                                                 image_size=(IMAGE_HEIGHT, IMAGE_WIDTH),
17                                                 color_mode="rgb",
18                                                 batch_size=1)
19
20     # get the true values and the predicted values in all the images from the test dataset
21     y_pred = []
22     y_true = []
23     for image, label in test_dataset:
24         y_pred.append(int(np.argmax(model.predict(image), axis=1)))
25         y_true.append(int(label))
26
27     sum = 0
28     for i in range(len(y_pred)):
29         if y_pred[i] == y_true[i]:
30             sum += 1
31
32     test_accuracy = sum / len(y_true)
33
34     cars_accuracy.update({vehicle: test_accuracy})
35
36     print('Test set accuracy for {}: {}'.format(vehicle, test_accuracy))
37     print(CLASSES.values)
38
39     # plot the confusion matrix
40     confusion_mtx = tf.math.confusion_matrix(y_true, y_pred)
41     plt.figure(figsize=(10, 8))
42     sns.heatmap(confusion_mtx, xticklabels=CLASSES, yticklabels=CLASSES,
43                 annot=True, fmt='g')
44     plt.xlabel('Prediction')
45     plt.ylabel('Label')
46     plt.show()
47
48 mean_accuracy = np.mean(list(cars_accuracy.values()))
49 print(cars_accuracy)
50 print("mean accuracy: {}".format(mean_accuracy))

```


Found 1500 files belonging to 7 classes.
 Test set accuracy for aclass: 0.544

<built-in method values of dict object at 0x7f4414d10870>




Found 1500 files belonging to 7 classes.
 Test set accuracy for escape: 0.43133333333333335


<built-in method values of dict object at 0x7f4414d10870>



Found 1500 files belonging to 7 classes.
 Test set accuracy for gsf: 0.43066666666666664

<built-in method values of dict object at 0x7f4414d10870>






Found 1010 files belonging to 7 classes.

Test set accuracy for i3: 0.5158415841584159


<built-in method values of dict object at 0x7f4414d10870>




Found 1500 files belonging to 7 classes.

Test set accuracy for tigan: 0.5346666666666666

<built-in method values of dict object at 0x7f4414d10870>






Found 1500 files belonging to 7 classes.

Test set accuracy for tucson: 0.5953333333333334


<built-in method values of dict object at 0x7f4414d10870>



Found 1500 files belonging to 7 classes.

Test set accuracy for x5: 0.6133333333333333

<built-in method values of dict object at 0x7f4414d10870>



{'aclass': 0.544, 'escape': 0.43133333333333335, 'gsf': 0.43066666666666664, 'i3': 0.5158415841584159, 'tigr': 0.5235964167845356}

▼ Testing on back seat images

```

1 import cv2
2 import matplotlib.pyplot as plt
3 import tensorflow as tf
4
5 SEAT_HEIGHT = 550
6 SEAT_WIDTH = 250
7 SEAT_SIZE = (SEAT_HEIGHT, SEAT_WIDTH)
8
9 Y_START = [70, 70, 70]
10 X_START = [130, 364, 582]
11
12
13 def crop_seats_from_image (image, y_start, x_start, seat_size):
14     cropped_seats = []
15     seat_height = seat_size[0]
16     seat_width = seat_size[1]
17     for x, y in zip(x_start, y_start):
18         cropped_seats.append(image[y:y+seat_height, x:x+seat_width])
19
20     return cropped_seats
21
22
23 def anotate_classes_on_image (image, class_texts, anotate_points):
24     for text, point in zip(class_texts, anotate_points):
25         print("")
26         cv2.putText(image, text, point, cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 255, 0), 2)
27
28     return image
29
30
31 def classify_seats (seat_images, model):
32     class_texts = []
33     class_labels = []
34
35     for seat_image in seat_images:
36
37         seat_image_resized = cv2.resize(seat_image, (IMAGE_WIDTH, IMAGE_HEIGHT), cv2.INTER_LINEAR)
38         seat_image_resized = tf.expand_dims(seat_image_resized, 0)
39
40         prediction = model.predict(seat_image_resized)
41         class_label = int(np.argmax(prediction, axis=1))
42
43         class_labels.append(class_label)
44         class_texts.append(CLASSES[class_label])
45
46     return class_labels, class_texts
47
48
49 # load images from vehicle
50
51 SAMPLES = 1
52 VEHICLE = "escape"
53 test_images_path = "{}\datasets\test\grayscale\{}\test_with_labels\grayscale_wholeImage".format(MAIN_PATH, VEH
54
55 # load images
56 image_files = [f for f in listdir(test_images_path) if
57                 path.isfile(path.join(test_images_path, f)) and f.endswith(".png")]
58
59 image_file_samples = random.sample(image_files, SAMPLES)
60
61 # Do the inference and show the seat classification
62 for image_file in image_file_samples:
63     image = cv2.imread(test_images_path + '/' + image_file)

```

```

64
65
66 cropped_seat_images = crop_seats_from_image(image, Y_START, X_START, SEAT_SIZE)
67
68 seat_labels, seat_class_texts = classify_seats(cropped_seat_images, model)
69
70 anotation_points = [(X_START[0],Y_START[0]),
71                     (X_START[1]+10,Y_START[1]),
72                     (X_START[2]+10,Y_START[2])]
73
74 anotetated_image = anotate_classes_on_image(image,seat_class_texts,anotation_points)
75
76 plt.figure(figsize=(12,12))
77 plt.imshow(anotetated_image)

```



Conclusion and Future Work

After training and evaluating the model we can see that the results obtained with the train dataset with different vehicles, are very similar to the average accuracy found by the SVIRO team in their benchmark on classifiers, around 50% when evaluated on unknown vehicles. One of the reasons why the results are low, for unknown vehicles, is that the model leans very easily on the texture of the objects in the train dataset and overfits. Another problem is that in the cropped image the objects appear to be very big, making the extraction of the relevant features harder. Besides that, looking at the confusion matrices we can see that there is a lot of misclassifications between "everyday objects" and "Empty Seats", probably because in many pictures, parts of objects and people from the adjacent seats appear on the image.

Nevertheless, this project was a great opportunity to understand the problem and the limits of the image classification approach. There are other solutions to explore in order to have a better vision of the objects, using the whole image, as object detection and segmentation.

✓ 4s completed at 1:07 AM

